

Master in Artificial Intelligence

Document
structure

Language
identification

Introduction to Human Language Technologies

1. Document structure



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona



Outline

Document
structure

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Outline

Document
structure

Searching textual
zones

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting
- 2 Language identification

Document types

Document
structure

Searching textual
zones

Language
identification

- Documents containing text:
 - Structured documents (e.g., web pages being tables)
 - Semi-structured documents (e.g., web pages containing pieces of plain text, figures and tables)
 - Documents with plain text only (e.g., text files, emails, tweets, oral transcripts)

Accessing to plain text contained in web pages may be relevant.

XML Parsers

- Transform an XML/HTML/XHTML document into a tree of standard objects.
- Provide an interface to manage that tree.
- Textual zones in the document can be extracted from that tree using the interface.

Document
structure

Searching textual
zones

Language
identification

```
<?xml version="1.0"?>
<doc type="novel" title="The green apple">
<chapter id="1">
<p>There are lots of trees in Amsteel Hill. I remember
going there and spend all the morning climbing those
trees, trying to get as many apples as possible.</p>
<p> James always wanted to come with me but he
was too young to get climbing.</p>
...
</doc>
```

Using ElementTree.py

```
for c in root:
    lp=c.findall('p')
    for p in lp:
        print p.text
```

Outline

Document
structure

Tokenization

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Goal of tokenization

- Goal: split plain text into *basic units*
- Use: IR tasks, text categorization, sentence splitting, language identification, text normalization . . .
- Different *basic units* depending on the task,
 - *Naïve* tokenizations: split by blanks and punctuation marks occurring after alphanum-string.
 - Complex tokenizations: names, clitics, abbreviations, **collocations**. . .

Goal of tokenization

- Goal: split plain text into *basic units*
- Use: IR tasks, text categorization, sentence splitting, language identification, text normalization ...
- Different *basic units* depending on the task,
 - Naïve tokenizations: split by blanks and punctuation marks occurring after alphanumeric-string.
 - Complex tokenizations: names, clitics, abbreviations, **collocations**...

Relevant definitions:

Word N-gram: sequence of words occurring in a text

Collocation: sequence of words that frequently occur together. Ex: "break a leg", "On the one hand"

Examples of tokenization

| Blanks | outer punct. | Abbr. | Clitics | Colloc. | text normalized |
|-----------|--------------|--------|---------|-----------|-----------------|
| Of | Of | Of | Of | Of_course | Of_course |
| course | course | course | course | | |
| I'll | I'll | I'll | I | I | I |
| | | | 'll | 'll | will |
| go | go | go | go | go | go |
| to | to | to | to | to | to |
| U.P.C. | U.P.C | U.P.C | U.P.C | U.P.C | Universitat... |
| | . | . | . | . | . |
| "Daily, | Daily | Daily | Daily | Daily | Daily |
| | , | , | , | , | , |
| Mr. | Mr | Mr. | Mr. | Mr. | Mister |
| | . | . | . | . | . |
| John | John | John | John | John | John_Smith |
| Smith..." | Smith | Smith | Smith | Smith | |
| | ... | ... | ... | ... | ... |
| | " | " | " | " | " |

Document
structure

Tokenization

Language
identification

Examples of tokenization

| Blanks | outer punct. | Abbr. | Clitics | Colloc. | text normalized |
|-----------|--------------|--------|---------|-----------|-----------------|
| Of | Of | Of | Of | Of_course | Of_course |
| course | course | course | course | | |
| I'll | I'll | I'll | I | I | I |
| | | | 'll | 'll | will |
| go | go | go | go | go | go |
| to | to | to | to | to | to |
| U.P.C. | U.P.C | U.P.C | U.P.C | U.P.C | Universitat... |
| | . | . | . | . | . |
| "Daily, | Daily | Daily | Daily | Daily | Daily |
| | , | , | , | , | , |
| Mr. | Mr | Mr. | Mr. | Mr. | Mister |
| | . | | | | |
| John | John | John | John | John | John_Smith |
| Smith..." | Smith | Smith | Smith | Smith | |
| | ... | ... | ... | ... | ... |
| | " | " | " | " | " |

Document
structure

Tokenization

Language
identification

Examples of tokenization

| Blanks | outer punct. | Abbr. | Clitics | Colloc. | text normalized |
|----------------------|----------------------|----------------------|--------------------------|---------------------------|----------------------------|
| Of course I'll | Of course I'll | Of course I'll | Of course I 'll | Of_course I 'll | Of_course I will |
| go to U.P.C. | go to U.P.C. | go to U.P.C | go to U.P.C | go to U.P.C | go to Universitat... |
| | . | . | . | . | . |
| "Daily, | Daily | Daily | Daily | Daily | Daily |
| Mr. | , | , | , | , | , |
| | Mr | Mr. | Mr. | Mr. | Mister |
| John Smith..." | . | . | . | . | . |
| | John Smith | John Smith | John Smith | John Smith | John_Smith |
| | ... | ... | ... | ... | ... |
| | " | " | " | " | " |

Problems: Non-standard text? Chinese? Japanese?

Document
structure
Tokenization
Language
identification

Outline

Document
structure

Sentence splitting

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Goal of sentence splitting

- Goal: Recognition of sentence boundaries in plain text (e.g., '. ' '? '!' '...').
- Language-dependent task
 - Ex: German: "Mein 2. Semester kommt bald zu Ende."
 - Ex: Traditional chinese?
- Domain-dependent task
 - Ex: "It is expressed as $(x=1)$? T.add('-') : T.add(x)."
- Methods:
 - Hand-crafted rules
 - Machine learning methods
- Input:
 - Naïve tokenization that depends on the particular method.
 - For simplicity, we will assume *blanks+outer_punctuation*
 - " I'll go to U.P.C. "Daily, Mr. John Smith..." "
 - " I 'll go to U.P.C . " Daily , Mr . John Smith ... " "

Problems of sentence splitting

Main problems:

- Abbreviations and acronyms (most difficult one)

Ex: "I will meet with Mr. Smith to talk about it."

Ex: "Lisa run 25 km. She ended up in N.Y."

How to detect them?

- Ellipsis

Ex: "There're different methods (A, B, ...) but ..."

- Internal quotation

Ex: " 'Stop!' he shouted."

- Ordinal numbers (German)

- Special cases:

Ex: " We have some variables. x stands for the weight,"

Hand-crafted rules for sentence splitting

- Specific hand-crafted rules for specific cases
 - Abbreviation classes (Lists of abbreviations)
(month name, unit-of-measure, title, address name, ...)
Ex: TITLE=('Mr', 'Mrs', 'Dr', ...)
 - Regular expressions for general cases, abbreviations, ellipsis, ...
Ex: / ([?!]+) / $\rightarrow t \in s_boundary$
Ex: / (\.\.\.\.) [A-Z]/ $\rightarrow t \in s_boundary$
Ex: / ([?!.]) [A-Z]/ $\rightarrow t \in s_boundary$
Ex: / (\$TITLE) \. / $\rightarrow t \notin s_boundary$
Ex: / [A-Z] \. / $\rightarrow t \notin s_boundary$
- Problem:
 - Highly expensive adaptation to new languages
(rules and abbreviation classes)

Supervised ML for sentence splitting

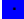

- The most frequently used (ME, SVM, Perceptron, ...-discriminative methods-)
- Requires manually annotated corpora. Commonly, $e^+, e^- = [',', '!', '?']$ and some preceding and following tokens
- Represents each e as a set of features, which depend on the approach, the language and the domain, although normally they tend to be binary features.
- Problem:
 - Requires very large sets of examples (tens of thousands to hundreds of thousands)

Supervised ML for sentence splitting

- Examples of features used in the state of the art
 - tok-1_X: 1st token before '.' is X
 - tok-2_X: 2nd token before '.' is X
 - tok+1_X: 1st token after '.' is X
 - len_tok-1_X: length of 1st token before '.' is X
 - len_tok-2_X: length of 2nd token before '.' is X
 - len_tok+1_X: length of 1st token after '.' is X
 - [up|lo|cap|num]_tok-1: 1st token before '.' is Upper, Lower, CAP, Numbers
 - [up|lo|cap|num]_tok-2: same for 2nd token before '.'
 - [up|lo|cap|num]_tok+1: same for 1st token after '.'
 - class_tok-1_X: abbreviation class of 1st token before '.' is X
 - ...

Supervised ML for sentence splitting

Example of annotation and binary features extraction

I 'll go to U.P.C  " Daily , Mr  John Smith ... "

| | | | |
|-------|-------------|-------|-------------------|
| e^+ | tok-1_U.P.C | e^- | tok-1_Mr |
| | len_tok-1_3 | | len_tok-1_2 |
| | CAP_tok-1 | | up_tok-1 |
| | tok-2_to | | tok-2, |
| | len_tok-2_2 | | len_tok-2_1 |
| | lo_tok-2 | | class_tok-1_title |
| | tok+1_" | | tok+1_John |
| | len_tok+1_1 | | len_tok+1_4 |
| | | | up_tok+1 |

Unsupervised ML for sentence splitting

- Based on corpus statistics
- Easily adaptable to new languages
 - They require large unannotated training corpora
- Mainly focus on abbreviations and ellipsis
- Heuristics and statistics calculated from the training corpus to decide:
 - 1 Which tokens are abbreviations?
 - 2 When the final period of the elements is a sentence boundary?
- Example: Punkt [Kiss and Strunk, 2006]

Unsupervised ML for sentence splitting

1 Punkt: Is token t considered an abbreviation?

Measured by considering the following heuristics:

- $t' = \langle t, . \rangle$ should be a collocation
- the length of t should be short
- t could include periods (acronyms)
- t is not ordinary word preceeding a period most of the times. (e.g., verbs in Turkish)

Unsupervised ML for sentence splitting

1 Punkt: Is token t considered an abbreviation?

Measured by considering the following heuristics:

- $t' = \langle t, . \rangle$ should be a collocation
- the length of t should be short
- t could include periods (acronyms)
- t is not ordinary word preceeding a period most of the times. (e.g., verbs in Turkish)

2 Punkt: Is the final period of abbreviation $t' = \langle t, . \rangle$ considered a sentence boundary?

Either one of the following heuristics must be true:

- $t'' = \text{following}(t')$ is a frequent sentence (from [1]) starter
- t'' is uppercase, occurs at least once in lowercase in the training corpus but never in uppercase inside sentences (from [1])

Exercise

Explain why Punkt fails (red) or not (blue) with the following texts:

- " "Good night!", said Laura. "
- " Abbrev. is a common abbreviation of abbreviation. "
- " We are meeting with our mr. You are late! "
- " We are meeting with our Mr. However, we'll finish soon."

Demo sentence splitters:

<http://text-processing.com/demo/tokenize/>

Outline

Document
structure

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Goal of language identification

Document
structure

Language
identification

- Can be seen as a particular classification problem.
- Given a document, d , and a set of languages, $L = \{l_1, \dots, l_k\}$, assign l_i to d .
- Method:
 - $\hat{d} = \text{representation}(d)$
 - $M(\hat{d}) \rightarrow l_i$
- Model M can be learned from training corpus $T = \{T_i\}_{1 \dots k}$ where $T_i = \{d_x | d_x \text{ written in } l_i\}$:
 - Supervised Machine Learning methods
 - Statistical Language models

Survey: <https://arxiv.org/pdf/1804.08186.pdf>

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

- 1 Which is the representation \hat{d} ?
- 2 How is $P^{T_i}(\hat{d})$ computed?

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

- 1 Which is the representation \hat{d} ?
- 2 How is $P^{T_i}(\hat{d})$ computed?

They depend on the particular type of model.

Most frequently used: **unigram language models**

Unigram language models for language identification

1 Which is the representation \hat{d} ?

$\hat{d} = e_1, \dots, e_s$ being the occurrences of unigrams:

- Words (after *Naïve* tokenization) or
- Characters n -grams (tokenization is not required)
 - n fixed (the most frequently used) or
 - n variable (improves accuracy, lower efficiency)

Unigram language models for language identification

1 Which is the representation \hat{d} ?

$\hat{d} = e_1, \dots, e_s$ being the occurrences of unigrams:

- Words (after *Naïve* tokenization) or
- Characters n -grams (tokenization is not required)
 - n fixed (the most frequently used) or
 - n variable (improves accuracy, lower efficiency)

2 How is $P^T(\hat{d})$ computed?

Each e_j is independent from the rest

$$P^T(\hat{d}) = P^T(e_1, \dots, e_s) = \prod_{j=1}^s P^T(e_j)$$

$$\log P^T(\hat{d}) = \sum_{j=1}^s \log P^T(e_j)$$

Possible estimators of $P^T(e_j)$:

- Maximum Likelihood Estimator (MLE)
- Smoothing techniques.

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

- Problem: data sparseness. Unseen e_j causes the model to fail. MLE is unsuitable for NLP.

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

■ Example:

$P^{[en]}('The\ doctor\ tell\ us\ about\ his\ quadriplegia')$?

$$c_{[en]}('quadriplegia') = 0 \implies P_{MLE}^{[en]}('quadriplegia') = 0$$

$$\implies P^{[en]}('The\ doctor\ tell\ us\ about\ his\ quadriplegia') = 0 !!$$

Unigram language models for language identification

Document
structure

Language
identification

Smoothing Techniques:

Keep some probability mass for e_j unseen in T_i

E.g., Lidstone's Law (LID)

$$P^T(e_j) \approx P_{LID}^T(e_j) = \frac{c_T(e_j) + \lambda}{N_T + \lambda B} \quad \text{usually, } \lambda = 0,5$$

B : #bins (potentially observable unigrams)

Exercise

Suppose we have a Language Identifier for English and Catalan, based on unigram language models with words and the following statistics

| w_i | a | he | mail | sent | to | mordorian |
|-----------------------------|-------------------|--------|-------|------|--------|-----------|
| English language model [en] | | | | | | |
| $c_{[en]}(w_i)$ | 17.000 | 10.000 | 3.900 | 850 | 25.000 | 0 |
| $N_{[en]}=1.300.000$ | $B_{[en]}=22.600$ | | | | | |
| Catalan Language model [ca] | | | | | | |
| $c_{[ca]}(w_i)$ | 21.000 | 11.900 | 420 | 910 | 750 | 0 |
| $N_{[ca]}=1.100.000$ | $B_{[ca]}=36.800$ | | | | | |

- Compute $P^{[en]}$ and $P^{[ca]}$ using MLE and LID for the following texts:
 - "he"
 - "he sent a"
 - "he sent a mail"
 - "he sent a mail to a mordorian"
- What language is identified by each estimator for each of the previous texts?
- Explain the effects of the text size