

# LECTURE 3b: Practical Reasoning Agents

Introduction to Multi-Agent Systems (MESIIA, MIA)

URV

# Classes of Architecture

- 1956–present: Symbolic Reasoning Agents
  - Agents make decisions about what to do via symbol manipulation.
  - Its purest expression, proposes that agents use explicit logical reasoning in order to decide what to do.
- 1985–present: Reactive Agents
  - Problems with symbolic reasoning led to a reaction against this
  - led to the reactive agents movement
- 1990-present: Hybrid Agents
  - Hybrid architectures attempt to combine the best of reasoning and reactive architectures

# What is Practical Reasoning?

- Practical reasoning is reasoning directed towards actions — the process of figuring out what to do:

## **Bratman (1990)**

“... Practical reasoning is a matter of weighing conflicting considerations for and against competing options, where the relevant considerations are provided by what the agent desires/values/cares about and what the agent believes...”

- Distinguish practical reasoning from *theoretical reasoning*
  - Theoretical reasoning is directed towards *beliefs*.

# The components of Practical Reasoning

- Human practical reasoning consists of two activities:
  - Deliberation:
    - Deciding *what* state of affairs we want to achieve
    - The output of deliberation is *intentions*;
  - Means-ends reasoning
    - Deciding *how* to achieve these state of affairs
    - The outputs of this step are *plans*
- Intentions are a key part of this.
  - The interaction between *beliefs*, *desires* and *intentions* defines how the model works.

# Intentions in Practical Reasoning

- Future-directed intentions
  - Intentions that an agent has towards some future state affairs.
- Intentions are:
  - *Pro-attitudes*, they tend to lead to action (drive means-ends reasoning).
  - *Persist*, attempt to achieve it (but not for too long).
  - *Consistent*/ constrain the future deliberation while holding some particular intentions, no options that are inconsistent with the intentions can be considered.
  - Influence beliefs upon with future practical reasoning is based.

# Intentions in Practical Reasoning

1. Intentions pose problems for agents, who need to determine ways of achieving them. Type equation here.

*If I have an intention to  $\phi$ , you would expect me to devote resources to deciding how to bring about  $\phi$*

2. Intentions provide a “filter” for adopting other intentions, which must not conflict.

*If I have an intention to  $\phi$ , you would not expect me to adopt an intention  $\psi$  that was incompatible with  $\phi$ .*

3. Agents track the success of their intentions, and tend to try again if their attempts fail.

*If an agent's first attempt to achieve  $\phi$  fails, then all other things being equal, it will try an alternative plan to achieve  $\phi$ .*

# Intentions in Practical Reasoning

4. Agents believe their intentions are possible.

*That is, they believe there is at least some way that the intentions could be brought about.*

5. Agents do not believe they will not bring about their intentions.

*It would not be rational of me to adopt an intention to  $\phi$  if I believed I would fail with  $\phi$ .*

6. Under certain circumstances, agents believe they will bring about their intentions.

*If I intend  $\phi$ , then I believe that under “normal circumstances” I will succeed with  $\phi$ .*

# Intentions in Practical Reasoning

7. Agents need not intend all the expected side effects of their intentions.

*If I believe  $\phi \Rightarrow \psi$  and I intend that  $\phi$ , I do not necessarily intend  $\psi$  also.*

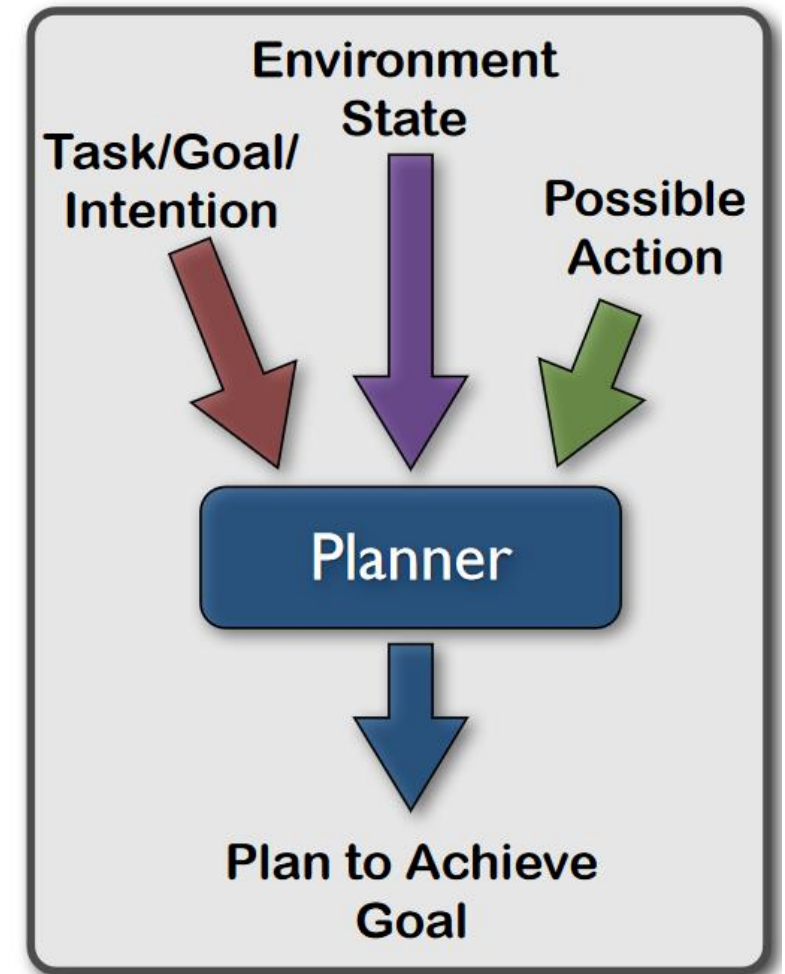
8. Intentions are not closed under implication.

*I may believe that going to the dentist involves pain, and I may also intend to go to the dentist — but this does not imply that I intend to suffer pain!*



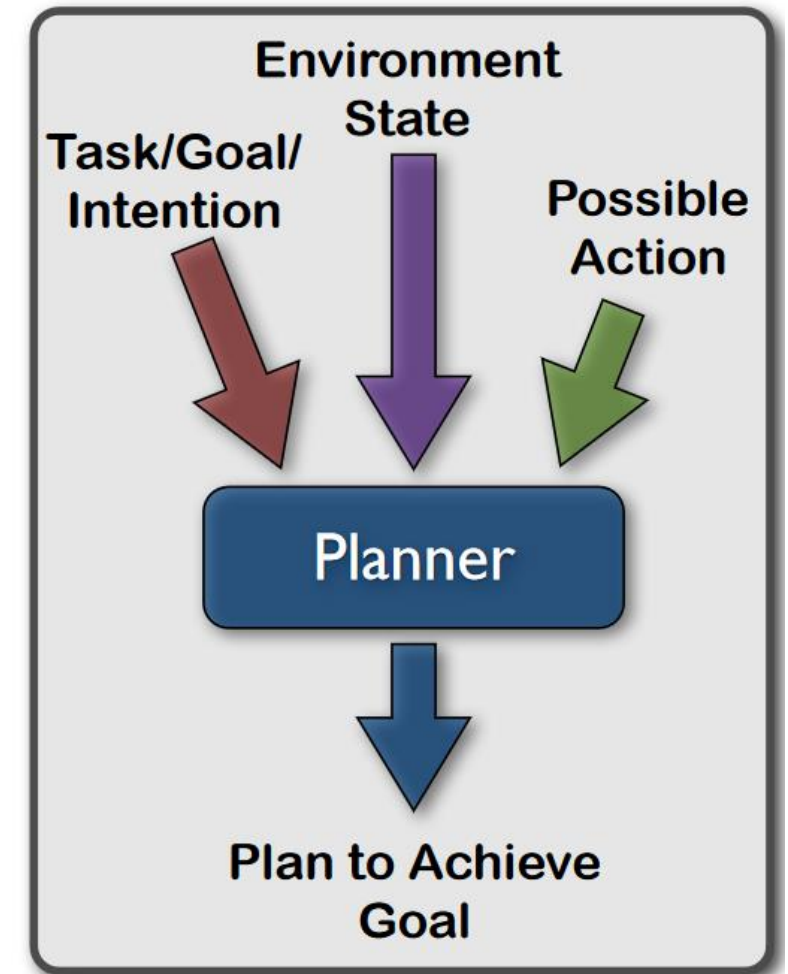
# Means-ends Reasoning/Planning

- Planning is the design of a course of action that will achieve some desired goal.
  - Basic idea is to give a planning system:
    - (representation of) goal/intention to achieve;
    - (representation of) actions it can perform;
    - (representation of) the environment.
  - and have it generate a *plan* to achieve the goal.
- This is ***automatic programming***.




# Means-ends Reasoning/Planning

- Don't have to directly tell the system what to do!
  - Let it *figure out* how to achieve the goal on its own.



# STRIPS Planner

- The **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- Used by Shakey, the robot developed by Richard Fikes and Nils Nilsson in 1971 at SRI International.




# Representations

- Question: How do we represent. . .
  - *goal to be achieved;*
  - *state of environment;*
  - *actions available to agent;*
  - *plan itself.*
- Answer: We use *logic*, or something that looks a lot like logic.

# Blocksworld

- We'll illustrate the techniques with reference to the blocks world.
  - A simple (toy) world, in this case one where we consider toys.
- The blocks world contains a robot arm, 3 blocks (A, B and C) of equal size, and a table-top.
- The aim is to generate a plan for the robot arm to build towers out of blocks.



# Blocksworld

- The environment is represented by an *ontology*.
- The closed world assumption is used
  - Anything not stated is assumed to be false.
- A *goal* is represented as a set of formulae

## Blocksworld Ontology

<i>On(x,y)</i>	object <i>x</i> on top of object <i>y</i>
<i>OnTable(x)</i>	object <i>x</i> is on the table
<i>Clear(x)</i>	nothing is on top of object <i>x</i>
<i>Holding(x)</i>	arm is holding <i>x</i>

## Representation of the following blocks

*Clear(A)*  
*On(A, B)*  
*OnTable(B)*  
*Clear(C)*  
*OnTable(C)*  
*ArmEmpty*

### The goal:

*{OnTable(A), OnTable(B), OnTable(C), ArmEmpty}*

# Blocksworld Actions

- Each action has:
  - a *name*: which may have arguments
  - a *pre-condition list*: list of facts which must be true for action to be executed;
  - a *delete list*: list of facts that are no longer true after action is performed;
  - an *add list*: list of facts made true by executing the action.
- Each of these may contain variables
- What is a plan?
  - A sequence (list) of actions, with *variables replaced by constants*.

# Blocksworld Actions

*Stack( $x, y$ )*  
pre *Clear( $y$ )  $\wedge$  Holding( $x$ )*  
del *Clear( $y$ )  $\wedge$  Holding( $x$ )*  
add *ArmEmpty  $\wedge$  On( $x, y$ )*

The **stack** action occurs when the robot arm places the object  $x$  it is holding is placed on top of object  $y$ .

*Pickup( $x$ )*  
pre *Clear( $x$ )  $\wedge$  OnTable( $x$ )  $\wedge$  ArmEmpty*  
del *OnTable( $x$ )  $\wedge$  ArmEmpty*  
add *Holding( $x$ )*

The **pickup** action occurs when the arm picks up an object  $x$  from the table.

*UnStack( $x, y$ )*  
pre *On( $x, y$ )  $\wedge$  Clear( $x$ )  $\wedge$  ArmEmpty*  
del *On( $x, y$ )  $\wedge$  ArmEmpty*  
add *Holding( $x$ )  $\wedge$  Clear( $y$ )*

The **unstack** action occurs when the robot arm picks an object  $y$  up from on top of another object  $y$ .


*PutDown( $x$ )*  
pre *Holding( $x$ )*  
del *Holding( $x$ )*  
add *OnTable( $x$ )  $\wedge$  ArmEmpty  $\wedge$  Clear( $x$ )*

The **putdown** action occurs when the arm places the object  $x$  onto the table.



# Using Plans

To get from here (left) to here (right)...



...we need this set of actions:

*UnStack(A,B)*

*Putdown(A)*

*Pickup(B)*

*Stack(B, C)*

*Pickup(A)*

*Stack(A, B)*

*Stack(x, y)*

pre *Clear(y) ∧ Holding(x)*

del *Clear(y) ∧ Holding(x)*

add *ArmEmpty ∧ On(x, y)*

*UnStack(x, y)*

pre *On(x, y) ∧ Clear(x) ∧ ArmEmpty*

del *On(x, y) ∧ ArmEmpty*

add *Holding(x) ∧ Clear(y)*

*Pickup(x)*

pre *Clear(x) ∧ OnTable(x) ∧ ArmEmpty*

del *OnTable(x) ∧ ArmEmpty*

add *Holding(x)*

*PutDown(x)*

pre *Holding(x)*

del *Holding(x)*

add *OnTable(x) ∧ ArmEmpty ∧ Clear(x)*

# Plan Validity

- Thus, a plan is simply a sequence of steps
- However, how can we:
  - Generate the plan?
  - Ensure that it is correct?

# Formal Representation

- As before we assume that the agent has a set of actions  $A_c$ , and we will write individual actions as  $\alpha_1, \alpha_2$  and so on.
- Now, we define a new component called action *descriptor*.

# Formal Representation

- A descriptor for an action  $\alpha \in Ac$  is a triple  $\langle P_\alpha, D_\alpha, A_\alpha \rangle$  where
  - $P_\alpha$  is the set of formulae of first-order logic that describes the *preconditions* of action  $\alpha$ .
  - $D_\alpha$  is the set of formulae of first-order logic that represents those facts will be *false* by applying the action  $\alpha$  (the *delete* list).
  - $A_\alpha$  is the set of formulae of first-order logic that represents those facts will be *true* by applying the action  $\alpha$  (the *add* list).
- For simplicity, we assume these lists constrained to only contain *ground atoms*.

# Formal Representation

- A *planning problem* is determined by a triple  $\langle \Delta, O, \gamma \rangle$  where
  - $\Delta$  is the beliefs of the agent about the *initial* state of the world.
  - $O = \{ \langle P_\alpha, D_\alpha, A_\alpha \rangle \mid \alpha \in Ac \}$  is an indexed set of operator descriptors, one for each available action  $\alpha$ .
  - $\gamma$  is a set of formulae of first-order logic, representing the *goal/task/intention* to be achieved.
- So, a *plan*  $\pi$  is a sequence of actions  $\pi = (\alpha_1, \alpha_2, \dots, \alpha_n)$  where  $\alpha_i$  is a member of  $Ac$ .

# Formal Representation

- With respect to a planning problem  $\langle \Delta, 0, \gamma \rangle$  a plan  $\pi = (\alpha_1, \alpha_2, \dots, \alpha_n)$  determines a sequence of  $n + 1$  belief databases  $\Delta_0, \Delta_1, \dots, \Delta_n$ , where  $\Delta_0 = \Delta$ , and

$$\Delta_i = (\Delta_{i-1} \setminus D_{\alpha_i}) \cup A_{\alpha_i} \text{ for } 1 \leq i \leq n$$

# Formal Representation

- Acceptable Plan:

A (linear) plan  $\pi = (\alpha_1, \alpha_2, \dots, \alpha_n)$  is said to be acceptable w.r.t the problem  $\langle \Delta, 0, \gamma \rangle$  iff the preconditions of every action is satisfied in preceding belief database, i.e., if  $\Delta_{i-1} \models P_{\alpha_i}$ , for all  $1 \leq i \leq n$ .

- Correct Plan:

A plan  $\pi = (\alpha_1, \alpha_2, \dots, \alpha_n)$  is correct w.r.t  $\langle \Delta, 0, \gamma \rangle$  iff:

1. It is acceptable
2.  $\Delta_n \models \gamma$  (i.e., if the goal is achieved in the final belief database generated by  $\pi$ )

# Action Definitions are important!!!

<b>Stack(x, y)</b>	
<b>pre</b>	<i>Clear(y) &amp; Holding(x)</i>
<b>del</b>	<i>Clear(y) &amp; Holding(x)</i>
<b>add</b>	<i>ArmEmpty &amp; On(x, y)</i>
<b>UnStack(x, y)</b>	
<b>pre</b>	<i>On(x, y) &amp; Clear(x) &amp; ArmEmpty</i>
<b>del</b>	<i>On(x, y) &amp; ArmEmpty</i>
<b>add</b>	<i>Holding(x) &amp; Clear(y)</i>
<b>Pickup(x)</b>	
<b>pre</b>	<i>Clear(x) &amp; OnTable(x) &amp; ArmEmpty</i>
<b>del</b>	<i>OnTable(x) &amp; ArmEmpty</i>
<b>add</b>	<i>Holding(x)</i>
<b>Release(x)</b>	
<b>pre</b>	<i>Holding(x)</i>
<b>del</b>	<i>Holding(x)</i>
<b>add</b>	<i>OnTable(x) &amp; ArmEmpty</i>

One of these definitions works. The other doesn't!

<b>Grap(x)</b>	
<b>pre</b>	<i>Clear(x) &amp; OnTable(x) &amp; ArmEmpty</i>
<b>del</b>	<i>OnTable(x) &amp; ArmEmpty</i>
<b>add</b>	<i>Holding(x)</i>
<b>Build(x, y)</b>	
<b>pre</b>	<i>Clear(y) &amp; Holding(x)</i>
<b>del</b>	<i>Clear(y) &amp; Holding(x)</i>
<b>add</b>	<i>ArmEmpty &amp; On(x, y)</i>
<b>Drop(y)</b>	
<b>pre</b>	<i>Holding(y)</i>
<b>del</b>	<i>Holding(y)</i>
<b>add</b>	<i>OnTable(y) &amp; ArmEmpty &amp; Clear(y)</i>
<b>Demolish(x, y)</b>	
<b>pre</b>	<i>On(x, y) &amp; Clear(y) &amp; ArmEmpty</i>
<b>del</b>	<i>On(x, y) &amp; ArmEmpty</i>
<b>add</b>	<i>Holding(x) &amp; Clear(y)</i>



# Action Definitions are important!!!

Stack(x, y)	
pre	<i>Clear(y) &amp; Holding(x)</i>
del	<i>Clear(y) &amp; Holding(x)</i>
add	<i>ArmEmpty &amp; On(x, y)</i>

UnStack(x, y)	
pre	<i>On(x, y) &amp; Clear(x) &amp; ArmEmpty</i>
del	<i>On(x, y) &amp; ArmEmpty</i>
add	<i>Holding(x) &amp; Clear(y)</i>

Pickup(x)	
pre	<i>Clear(x) &amp; OnTable(x) &amp; ArmEmpty</i>
del	<i>OnTable(x) &amp; ArmEmpty</i>
add	<i>Holding(x)</i>

Release(x)	
pre	<i>Holding(x)</i>
del	<i>Holding(x)</i>
add	<i>OnTable(x) &amp; ArmEmpty &amp; Clear(x)</i>

One of these definitions works. The other doesn't!

Grap(x)	
pre	<i>Clear(x) &amp; OnTable(x) &amp; ArmEmpty</i>
del	<i>OnTable(x) &amp; ArmEmpty</i>
add	<i>Holding(x)</i>

Build(x, y)	
pre	<i>Clear(y) &amp; Holding(x)</i>
del	<i>Clear(y) &amp; Holding(x)</i>
add	<i>ArmEmpty &amp; On(x, y)</i>

Drop(y)	
pre	<i>Holding(y)</i>
del	<i>Holding(y)</i>
add	<i>OnTable(y) &amp; ArmEmpty &amp; Clear(y)</i>

Demolish(x, y)	
pre	<i>On(x, y) &amp; Clear(y) &amp; ArmEmpty</i>
del	<i>On(x, y) &amp; ArmEmpty</i>
add	<i>Holding(x) &amp; Clear(y)</i>

# Implementing Practical Reasoning Agents

- A first pass at an implementation of a practical reasoning agent:
- For now, we will not be concerned with stages 2 or 3.
  - These are related to the functions *see* and *next*

## Agent Control Loop Version 1


```
1. while true
2.     observe the world;
3.     update internal world model;
4.     deliberate about what intention
       to achieve next;
5.     use means-ends reasoning to get
       a plan for the intention;
6.     execute the plan
7. end while
```

# Implementing Practical Reasoning Agents

- see is as before  
 $see: E \rightarrow Per$
- Instead of the function next...
  - which takes a percept and used it to update the internal state of an agent
- ...we have a belief revision function:

$$brf: 2^{Bel} \times Per \rightarrow 2^{Bel}$$

Where,  $Bel$  is the set of all possible beliefs that an agent might have



# Implementing Practical Reasoning Agents

- Problem:
  - deliberation and means-ends reasoning processes are not *instantaneous*.
  - They have a *time cost*.
- Suppose that deliberation is *optimal*
  - The agent selects the optimal intention to achieve, then this is the best thing for the agent.
  - i.e., it maximises the expected utility.
- So, the agent selects an intention to achieve that would have been optimal at the time it observed the world.
  - This is calculative rationality.
- The *world may change* in the meantime.
  - Even if the agent can compute the right thing to do, it may not do the right thing.
  - Optimality is hard.

# Implementing Practical Reasoning Agents

- Let's make the algorithm more formal
  - The term  $I \subseteq Int$ , the set of intentions
  - Plan() is the planning function,
  - brf() is the belief revision function
  - and execute() is a function that executes each action in a plan.
- How might we implement these functions?

```
Agent Control Loop Version 2
1.   $B := B_0$ ; /* initial beliefs */
2.  while true do
3.      get next percept  $\rho$ ;
4.       $B := brf(B, \rho)$ ;
5.       $I := deliberate(B)$ ;
6.       $\pi := plan(B, I)$ ;
7.      execute( $\pi$ )
8.  end while
```

# Deliberation

- How does an agent *deliberate*?
  - begin by trying to understand what the *options* available to it are;
  - *choose* between them, and *commit* to some.
- Chosen options are then *intentions*.
- The *deliberate* function can be decomposed into two distinct functional components:
  - ***option generation***; and
  - ***filtering***

# Option Generation and Filtering

- Option Generation

- The agent generates a set of possible alternatives
- Represent option generation via a function, **options()**, which takes the agent's current beliefs and current intentions, and from them determines a set of options
- **desires**

$options : \mathcal{P}(Bel) \times \mathcal{P}(Int) \rightarrow \mathcal{P}(Des)$

- Filtering

- the agent chooses between competing alternatives, and commits to achieving them.
- In order to select between competing options, an agent uses a **filter()** function.
- **Intentions**

$filter : \mathcal{P}(Bel) \times \mathcal{P}(Des) \times \mathcal{P}(Int) \rightarrow \mathcal{P}(Int)$

# Implementing Practical Reasoning Agents

## Agent Control Loop Version 3

```
1.   $B := B_0$ ;  
2.   $I := I_0$ ;  
3.  while true do  
4.      get next percept  $\rho$ ;  
5.       $B := brf(B, \rho)$ ;  
6.       $D := options(B, I)$ ;  
7.       $I := filter(B, D, I)$ ;  
8.       $\pi := plan(B, I)$ ;  
9.       $execute(\pi)$   
10. end while
```



# Degrees of Commitment

- Blind commitment
  - A blindly committed agent will continue to maintain an intention until it believes the intention has actually been achieved. Blind commitment is also sometimes referred to as *fanatical* commitment.
- Single-minded commitment
  - A single-minded agent will continue to maintain an intention *until* it believes that *either* the intention has been achieved, *or else that it is no longer possible to achieve* the intention.
- Open-minded commitment
  - An open-minded agent will maintain an intention *as long as it is still believed* possible.

# Degrees of Commitment

- An agent has commitment both to:
  - *ends* (i.e., the state of affairs it wishes to bring about), and
  - *means* (i.e., the mechanism via which the agent wishes to achieve the state of affairs).
- Currently, our agent control loop is *overcommitted*, both to means and ends.
  - Modification: *replan* if ever a plan goes wrong.
  - However, to write the algorithm down we *need to refine* our notion of plan execution.

# Degrees of Commitment

- The previous version was *blindly committed to its means and its ends*
- If  $\pi$  is a plan, then:
  - $empty(\pi)$  is true if there are no more actions in the plan.
  - $hd(\pi)$  returns the first action in the plan.
  - $tail(\pi)$  returns the plan minus the head of the plan.
  - $sound(\pi, I, B)$  means that  $\pi$  is sound plan for  $I$  given  $B$ .

## Agent Control Loop Version 4

```
1.   $B := B_0$ ;  
2.   $I := I_0$ ;  
3.  while true do  
4.      get next percept  $\rho$ ;  
5.       $B := brf(B, \rho)$ ;  
6.       $D := options(B, I)$ ;  
7.       $I := filter(B, D, I)$ ;  
8.       $\pi := plan(B, I)$ ;  
9.      while not  $empty(\pi)$  do  
10.          $\alpha := hd(\pi)$ ;  
11.          $execute(\alpha)$ ;  
12.          $\pi := tail(\pi)$ ;  
13.         get next percept  $\rho$ ;  
14.          $B := brf(B, \rho)$ ;  
15.         if not  $sound(\pi, I, B)$  then  
16.              $\pi := plan(B, I)$   
17.         end-if  
18.     end-while  
19. end-while
```

# Degrees of Commitment

- Makes the control loop more reactive, able to change intention when the world changes.
  - i.e. it is not committed to its means (line 16)
- Still overcommitted to intentions (ends).
  - Never stops to consider whether or not its intentions are appropriate.

## Agent Control Loop Version 4

```
1.   $B := B_0$ ;  
2.   $I := I_0$ ;  
3.  while true do  
4.      get next percept  $\rho$ ;  
5.       $B := brf(B, \rho)$ ;  
6.       $D := options(B, I)$ ;  
7.       $I := filter(B, D, I)$ ;  
8.       $\pi := plan(B, I)$ ;  
9.      while not empty( $\pi$ ) do  
10.          $\alpha := hd(\pi)$ ;  
11.         execute( $\alpha$ );  
12.          $\pi := tail(\pi)$ ;  
13.         get next percept  $\rho$ ;  
14.          $B := brf(B, \rho)$ ;  
15.         if not sound( $\pi, I, B$ ) then  
16.              $\pi := plan(B, I)$   
17.         end-if  
18.     end-while  
19. end-while
```

# Single Minded Commitment

- Modification:
  - stop to determine whether intentions have succeeded or whether they are impossible
- Our agent now gets to reconsider its intentions once every time around the outer control loop (line 9), i.e., after:
  - it has completely executed a plan to achieve its current intentions; or
  - it believes it has achieved its current intentions; or
  - it believes its current intentions are no longer possible

## Agent Control Loop Version 5

```
1.   $B := B_0$ ;
2.   $I := I_0$ ;
3.  while true do
4.      get next percept  $\rho$ ;
5.       $B := brf(B, \rho)$ ;
6.       $D := options(B, I)$ ;
7.       $I := filter(B, D, I)$ ;
8.       $\pi := plan(B, I)$ ;
9.      while not( $empty(\pi)$ 
                or  $succeeded(I, B)$ 
                or  $impossible(I, B)$ ) do
10.          $\alpha := hd(\pi)$ ;
11.          $execute(\alpha)$ ;
12.          $\pi := tail(\pi)$ ;
13.         get next percept  $\rho$ ;
14.          $B := brf(B, \rho)$ ;
15.         if not  $sound(\pi, I, B)$  then
16.              $\pi := plan(B, I)$ 
17.         end-if
18.     end-while
19. end-while
```

# Open Minded Commitment

- In the previous version, our agent reconsiders its intentions once every time around the outer control loop.
- In this new version, our agent also reconsiders its intentions after every action (lines 15 & 16)
- But this intention reconsideration is **costly!**

## Agent Control Loop Version 6

```
1.   $B := B_0$ ;  
2.   $I := I_0$ ;  
3.  while true do  
4.      get next percept  $\rho$ ;  
5.       $B := brf(B, \rho)$ ;  
6.       $D := options(B, I)$ ;  
7.       $I := filter(B, D, I)$ ;  
8.       $\pi := plan(B, I)$ ;  
9.      while not ( $empty(\pi)$   
                or  $succeeded(I, B)$   
                or  $impossible(I, B)$ ) do  
10.          $\alpha := hd(\pi)$ ;  
11.          $execute(\alpha)$ ;  
12.          $\pi := tail(\pi)$ ;  
13.         get next percept  $\rho$ ;  
14.          $B := brf(B, \rho)$ ;  
15.          $D := options(B, I)$ ;  
16.          $I := filter(B, D, I)$ ;  
17.         if not  $sound(\pi, I, B)$  then  
18.              $\pi := plan(B, I)$   
19.         end-if  
20.     end-while  
21. end-while
```

# Intention Reconsideration

- A dilemma:
  - an agent that *does not stop* to reconsider its intentions sufficiently often *will continue to attempt* to achieve its intentions even after it is clear that they cannot be achieved, or that there is no longer any reason for achieving them;
  - an agent that *constantly* reconsiders its intentions may *spend insufficient time* actually working to achieve them, and hence runs the risk of never actually achieving them.
- Solution: incorporate an explicit meta-level control component, that decides whether or not to *reconsider*.

## Agent Control Loop Version 7

```
1.   $B := B_0$ ;  
2.   $I := I_0$ ;  
3.  while true do  
4.      get next percept  $\rho$ ;  
5.       $B := brf(B, \rho)$ ;  
6.       $D := options(B, I)$ ;  
7.       $I := filter(B, D, I)$ ;  
8.       $\pi := plan(B, I)$ ;  
9.      while not( $empty(\pi)$   
              or  $succeeded(I, B)$   
              or  $impossible(I, B)$ ) do  
10.          $\alpha := hd(\pi)$ ;  
11.          $execute(\alpha)$ ;  
12.          $\pi := tail(\pi)$ ;  
13.         get next percept  $\rho$ ;  
14.          $B := brf(B, \rho)$ ;  
15.         if  $reconsider(I, B)$  then  
16.              $D := options(B, I)$ ;  
17.              $I := filter(B, D, I)$ ;  
18.         end-if  
19.         if not  $sound(\pi, I, B)$  then  
20.              $\pi := plan(B, I)$   
21.         end-if  
22.     end-while  
23. end-while
```

# Intention Reconsideration

- The possible interactions between meta-level control and deliberation are:

Situation number	Chose to deliberate?	Changed intentions?	Would have changed intentions?	<i>reconsider(...)</i> optimal?
1	No	—	No	Yes
2	No	—	Yes	No
3	Yes	No	—	No
4	Yes	Yes	—	Yes

- An important assumption: cost of **reconsider(...)** is ***much*** less than the cost of the deliberation process itself.

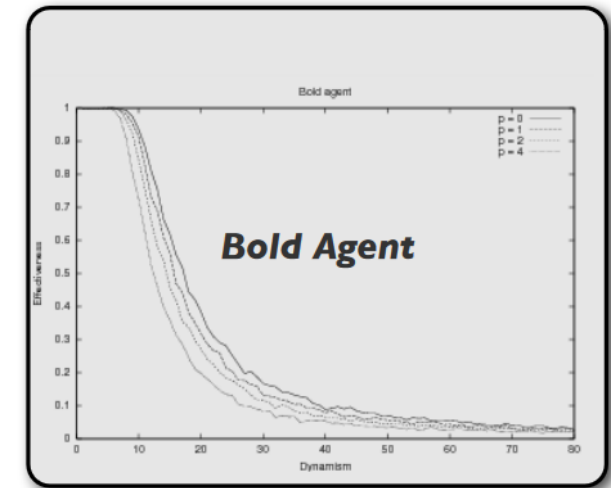


# Optimal Intention Reconsideration


- Two different types of reconsideration strategy considered:
  - bold agents: never pause to reconsider intentions, and
  - cautious agents: stop to reconsider after every action.
- Dynamism in the environment is represented by the rate of world change,  $\gamma$ .

# Optimal Intention Reconsideration

- if  $\gamma$  is low (i.e., the environment does not change quickly), then bold agents do well compared to cautious ones.
  - This is because cautious ones waste time reconsidering their commitments while bold agents are busy working towards — and achieving — their intentions.
- If  $\gamma$  is high (i.e., the environment changes frequently), then cautious agents can outperform bold agents.
  - This is because they are able to recognise when intentions are doomed, and also to take advantage of serendipitous situations and new opportunities when they arise.
- When planning costs are high, this advantage can be eroded.



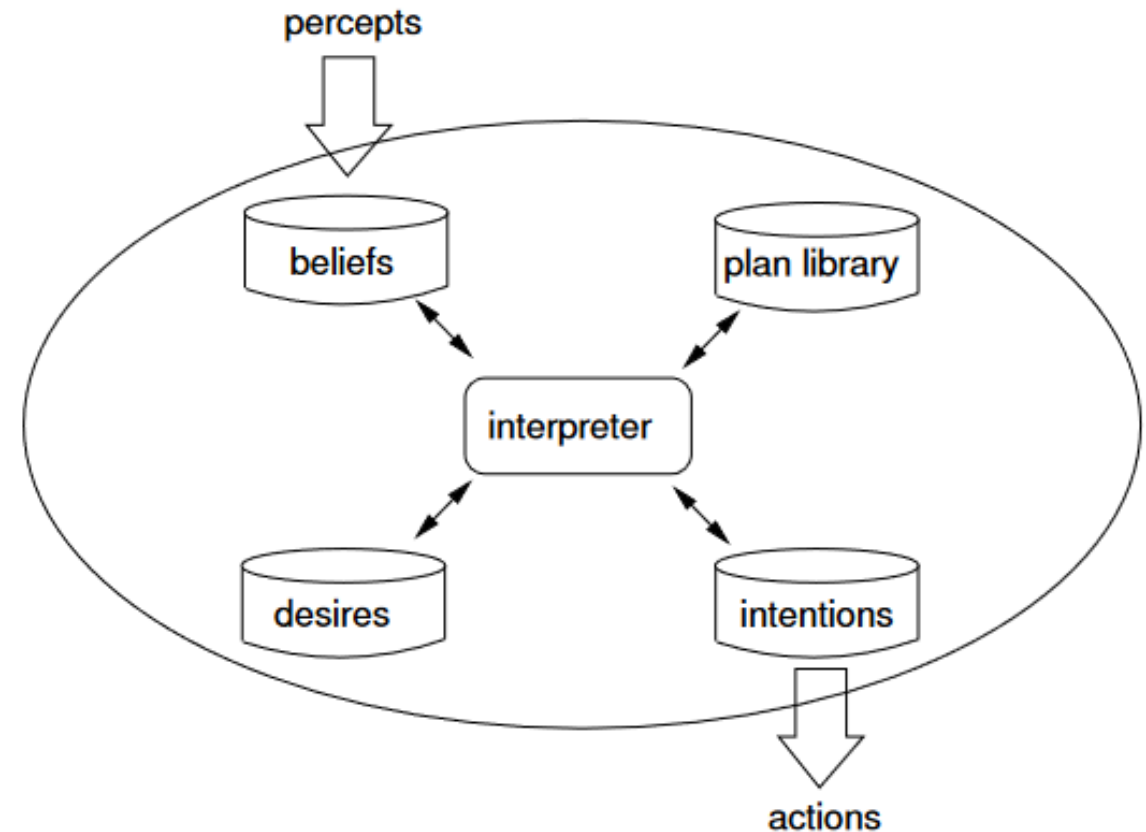
low  $\longleftarrow \gamma \longrightarrow$  high



# Implemented BDI Agents: Procedural Reasoning System (PRS)

---

- In the PRS, each agent is equipped with a plan library, representing that agent's procedural knowledge: knowledge about the mechanisms that can be used by the agent in order to realise its intentions.
- The options available to an agent are directly determined by the plans an agent has: an agent with no plans has no options.



# Readings for this week

- M.Wooldridge: An introduction to MultiAgent Systems – Ch. 4  
Practical Reasoning Agents