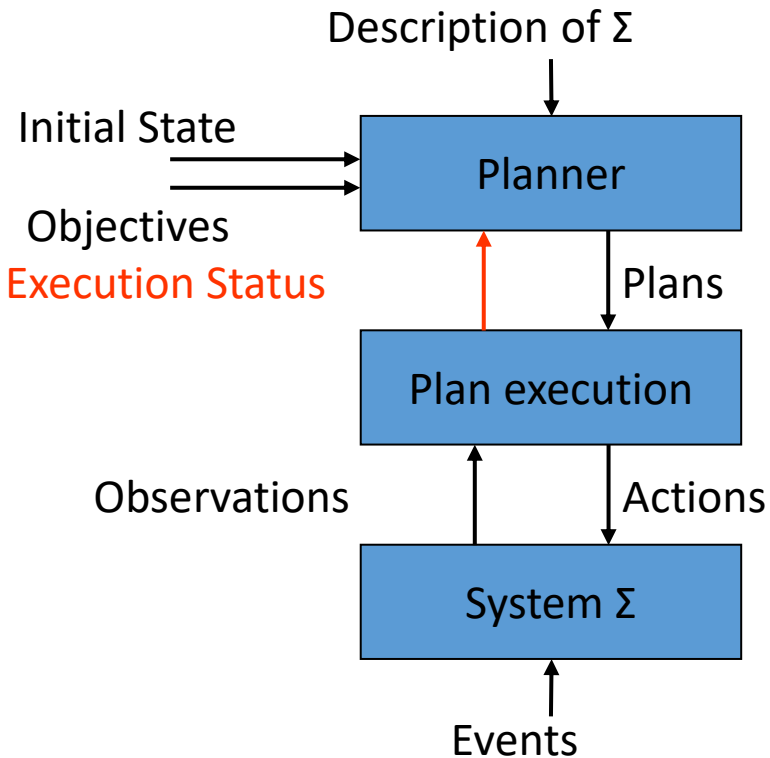


# Planning and Approximate Reasoning

Hatem A. Rashwan

## Planning under uncertainty: MDPs

# Plan Execution





- **problem**: real world differs from model described by  $\Sigma$
- **more realistic model**: interleaved planning and execution
  - plan supervision
  - plan revision
  - re-planning
- **dynamic planning**: closed loop between planner and controller
  - execution status

# Planning Example

## Robot navigation

- A robot act in the environment aiming to avoid the obstacles and reach its goal.
- The robot aims to reach the celled labelled with +1 and avoid the cell labelled with -1




# Planning Example

## Why the short path is problematic

- Non-perfect execution: actions are performed with a probability  $< 1$
- What are the best actions for a robot (an agent) under this constraint? (move right->up->up->right)
- And what is the maximum performance of the execution? (e.g., 99%)

### Example

- A robot does not exactly perform the desired motions due to different reasons.
  - Uncertainty about performing actions




# Planning Example

## Non-Deterministic Transitions

- Consider non-deterministic transition model (UP/R/L)

Desired action is UP





### Example

- Intended action is executed with  $p=0.8$
- With  $p = 0.1$ , the robot can move right or left.
  - Uncertainty about performing actions

# Markov Decision Process (MDP) motivation

- Executing the A\* plan




## Transitions are non-deterministic

- Uncertainty about performing actions will be occurred

# MDP motivation

- Perhaps using longer path with lower probability to not reach the cell labelled -1 is good option.



**This proposed path can have the highest overall utility.**

# Axioms of Probability

- Let  $A$  be a proposition about the world
- $P(A)$  = probability proposition  $A$  is true
- $0 \leq P(A) \leq 1$
- $P(\text{True}) = 1$
- $P(\text{False}) = 0$
- $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$


## Random Variables

- Random Variables: variables in probability to capture phenomena
- A random variable has a **domain of values** it can take on.
- Probability distribution function represents  
“**probability of each value**”



# Example – Pick Fruit from Basket

- Random variable:  $F$
- Domain:  $a, o$
- PDF:
  - $p(F = a) = \frac{1}{4}$
  - $p(F = o) = \frac{3}{4}$



- The **expected value** of a function of a random variable **is the weighted average of the probability distribution over outcomes.**
- Example: calculate the expected time of waiting for an elevator
- Time:            5ms    2ms    0.5ms
- Probability:    0.2    0.7    0.1
- $5 \times 0.2 + 2 \times 0.7 + 0.5 \times 0.1 = \mathbf{2.45ms}$

# Transition Model

- Given that our agent is in some state  $s$ , there is a probability to go to the first state, another probability to go to the second state, and so on for every existing state.
- This is our transition probability.
- Probability to reach the next state  $s'$  from state  $s$  by choosing action  $a$  is  $P(s, a, s') \sim P(s/s', a) \Rightarrow$  It is called **Transition model**

## Markov Property:

The transition probabilities from  $s$  to  $s'$  depend only on the current state  $s$  and not on the history of earlier states.

## Reward:

- In each state  $s$ , the robot (agent) receives a reward  $R(s)$ .
- The reward may be positive or negative but must be bounded
- This can be generalized to be a reward function  $R(s, a, s')$

# Reward

- In our example, the reward is **-0.04** in all states (e.g., **the cost of the motion**) except the terminal states (**reward is +1/-1**)

A negative reward gives an incentive

- to reach the goal quickly
- Or “living in this environment is not enjoyable”


-0.04	-0.04	-0.04	+1
-0.04		-0.04	-1
-0.04	-0.04	-0.04	-0.04

# Markov Decision Process-MDP


Given a **sequential decision problem** in **fully observable**, stochastic environment with a known Markovian transition model

Then a Markov Decision Process-MDP  $(S, A, S_i, P, R)$  is defined by the components of:

- **Set of states:**  $S$
- **Set of actions:**  $A$
- **Initial states:**  $S_i$
- **Transition model:**  $P(s, a, s')$
- **Reward function:**  $R(s)$ 
  - Here: considering only  $R(s)$   
(does not change the problem)



# Example – Markov System with Reward



- Process/observation:
  - Assume start state  $s_i$
  - Receive immediate reward  $r_i$
  - Move, or observe a move, randomly to a new state according to the probability transition matrix
  - **Future rewards (of next state) are discounted by  $\gamma$**

probability transition matrix  $T$  is

sun	wind	hail
1/2	1/2	0
1/2	0	1/2
0	1/2	1/2

# Example – Markov System with Reward


## Solving a Markov System with Rewards

- $V^*(s_i)$  - expected discounted sum of future rewards starting in state  $s_i$
- $V^*(s_i) = r_i + \gamma[p_{i1}V^*(s_1) + p_{i2}V^*(s_2) + \dots p_{in}V^*(s_n)]$
- $\gamma$  is a **discount factor**, where  $\gamma \in [0, 1]$ .
- It informs the agent of how much it should care about rewards now to rewards in the future.
- If ( $\gamma = 0$ ), that means the agent is **short-sighted**, in other words, it only cares about the first reward.
- If ( $\gamma = 1$ ), that means the agent is **far-sighted**, i.e. it cares about all future rewards.
- What we care about is the total rewards that we're going to get.

# Value Iteration to Solve a Markov System with Rewards

- $V_1(s_i)$  - expected discounted sum of future rewards starting in state  $s_i$  *for one step*.
- $V_2(s_i)$  - expected discounted sum of future rewards starting in state  $s_i$  *for two steps*.
- ...
- $V_k(s_i)$  - expected discounted sum of future rewards starting in state  $s_i$  *for  $k$  steps*.
- As  $k \rightarrow \infty$   $V_k(s_i) \rightarrow V^*(s_i)$
- Stop when difference of  $k + 1$  and  $k$  values is smaller than some  $\epsilon$ .

# Example – Markov System with Reward




## 3-State Example: Values $\gamma = 0.5$

Iteration	SUN	WIND	HAIL
0	0	0	0
1	4	0	-8
2	5.0	-1.0	-10.0
3	5.0	-1.25	-10.75
4	4.9375	-1.4375	-11.0
5	4.875	-1.515625	-11.109375
6	4.8398437	-1.5585937	-11.15625
7	4.8203125	-1.5791016	-11.178711
8	4.8103027	-1.5895996	-11.189453
9	4.805176	-1.5947876	-11.194763
10	4.802597	-1.5973969	-11.197388
11	4.8013	-1.5986977	-11.198696
12	4.8006506	-1.599349	-11.199348
13	4.8003254	-1.5996745	-11.199675
14	4.800163	-1.5998373	-11.199837
15	4.8000813	-1.5999185	-11.199919




# Example – Markov System with Reward



## 3-State Example: Values $\gamma = 0.9$

Iteration	SUN	WIND	HAIL
0	0	0	0
1	4	0	-8
2	5.8	-1.8	-11.6
3	5.8	-2.6100001	-14.030001
4	5.4355	-3.7035	-15.488001
5	4.7794	-4.5236254	-16.636175
6	4.1150985	-5.335549	-17.521912
7	3.4507973	-6.0330653	-18.285858
8	2.8379793	-6.6757774	-18.943516
9	2.272991	-7.247492	-19.528683
...	...	...	...
50	-2.8152928	-12.345073	-24.633476
51	-2.8221645	-12.351946	-24.640347
52	-2.8283496	-12.3581295	-24.646532
...	...	...	...
86	-2.882461	-12.412242	-24.700644
87	-2.882616	-12.412397	-24.700798
88	-2.8827558	-12.412536	-24.70094

# Example – Markov System with Reward



## 3-State Example: Values $\gamma = 0.2$



Iteration	SUN	WIND	HAIL
0	0	0	0
1	4	0	-8
2	4.4	-0.4	-8.8
3	4.4	-0.44000003	-8.92
4	4.396	-0.452	-8.936
5	4.3944	-0.454	-8.9388
6	4.39404	-0.45443997	-8.93928
7	4.39396	-0.45452395	-8.939372
8	4.393944	-0.4545412	-8.939389
9	4.3939404	-0.45454454	-8.939393
10	4.3939395	-0.45454526	-8.939394
11	4.3939395	-0.45454547	-8.939394
12	4.3939395	-0.45454547	-8.939394

# Markov Chain - Example

- Markov Chain
  - states
  - transitions
  - rewards
  - no actions
- Value of a state, using infinite discounted horizon
 
$$V^*(s_i) = r_i + \gamma[p_{i1}V^*(s_1) + p_{i2}V^*(s_2) + \dots p_{in}V^*(s_n)]$$
  - Assume  $\gamma=0.9$
$$V(1) = 0 + .9(.5 V(1) + .5 V(2))$$

$$V(2) = 10 + .9(.2 V(1) + .1 V(2) + .7 V(3))$$

$$V(3) = 0 + .9(.9 V(2) + .1 V(3))$$



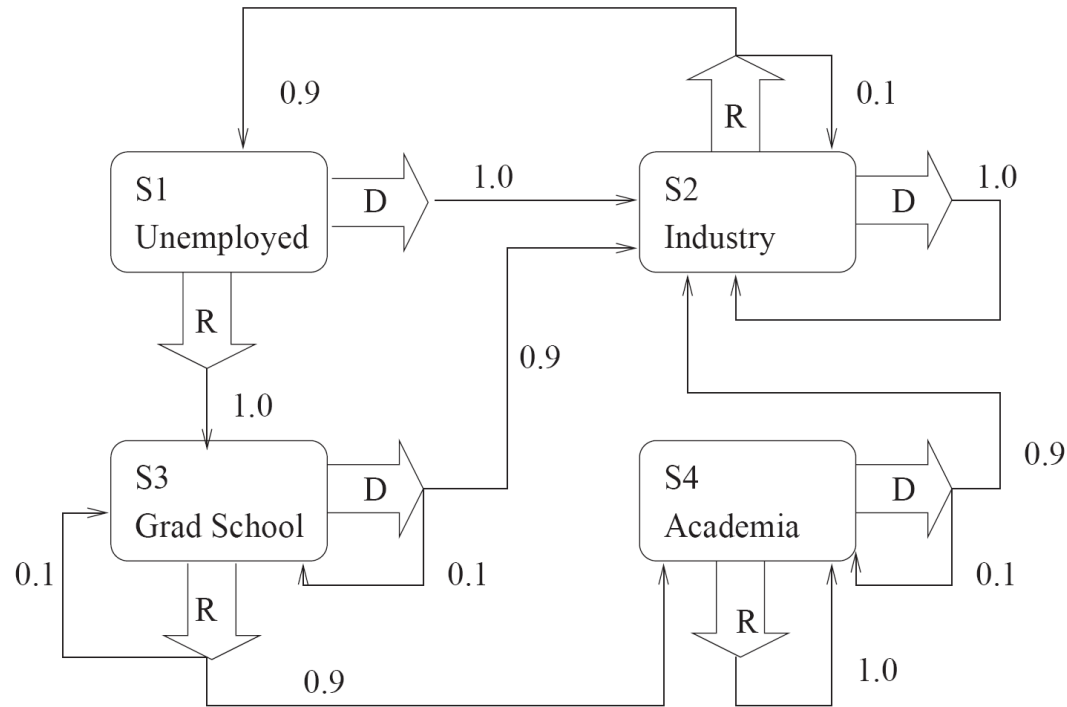
# Markov Decision Processes

- Finite set of states,  $s_1, \dots, s_n$
- **Finite set of actions,  $a_1, \dots, a_m$**
- Probabilistic state, action transitions:
- $P_{ij}^k = \text{prob}(\text{next} = s_j \mid \text{current} = s_i \mid \text{take action} = a_k)$
- **Markov property:** State transition function only dependent on current state, not on the “history” of how the state was reached.
- Reward for each state,  $r_1, \dots, r_n$
- Process:
  - Start in state  $s_i$
  - Receive immediate reward  $r_i$
  - **Choose action  $a_k \in A$**
  - Change to state  $s_j$  with probability  $P_{ij}^k$
  - Discount future rewards

# Solving an MDP

- Find an action to apply to each state.
- A **policy** is a mapping from states to actions.
- Optimal policy - for every state, **there is no other action that gets a higher sum of discounted future rewards.**
- For every MDP there exists an **optimal** policy.
- Solving an MDP is **finding an optimal policy.**
- A specific policy converts an MDP into a plain Markov system with rewards.

# Solving an MDP - Example



- Note the need to have a finite set of states and actions (R or D).  
**R=> Research, and D=> Development**
- Note the need to have all transition probabilities.


# Value Iteration

- $V^*(s_i)$  - expected discounted future rewards, if we start from state  $s_i$ , and we follow the optimal policy.
- Compute  $V^*$  with value iteration:
  - $V_k(s_i)$  = maximum possible future sum of rewards starting from state  $s_i$  for  $k$  steps.
- Bellman's Equation:

$$V^{n+1}(s_i) = \max_k \left\{ r_i + \gamma \sum_{j=1}^N P_{ij}^k V^n(s_j) \right\}$$

- Dynamic programming

# Nondeterministic Example



- Reward and discount factor to be decided.
- Note the need to have a finite set of states and actions.
- Note the need to have all transition probabilities.



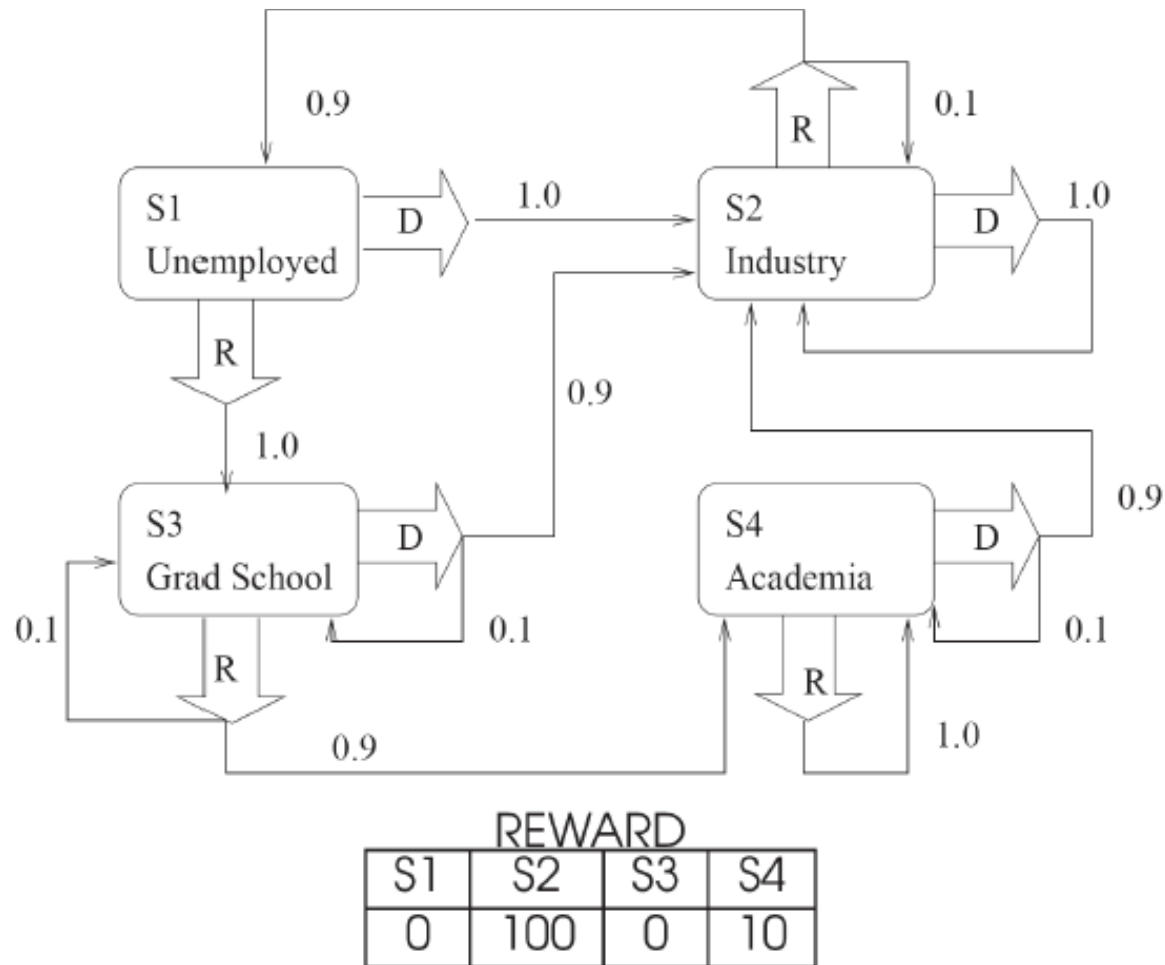
# Value Iteration-Bellman equation

- Start with some policy  $\pi_0(s_i)$ .
- Such policy transforms the MDP into a plain Markov system with rewards.
- Compute the values of the states according to the current policy.
- Update policy:

$$\pi_{k+1}(s_i) = r_i + \arg \max_a \left\{ \gamma \sum_{j=1}^N P_{ij}^a V^{\pi_n}(s_j) \right\}$$

- Keep computing
- Stop when  $\pi_{k+1} = \pi_k$ .

# Nondeterministic Example



# Nondeterministic Example

$\pi^*(s) = D$ , for any  $s = S1, S2, S3$ , and  $S4$ ,  $\gamma = 0.9$ .

---

$$V^*(S2) = r(S2, D) + 0.9 (1.0 V^*(S2))$$

$$V^*(S2) = 100 + 0.9 V^*(S2)$$

$$V^*(S2) = 1000.$$

$$V^*(S1) = r(S1, D) + 0.9 (1.0 V^*(S2))$$

$$V^*(S1) = 0 + 0.9 \times 1000$$

$$V^*(S1) = 900.$$

$$V^*(S3) = r(S3, D) + 0.9 (0.9 V^*(S2) + 0.1 V^*(S3))$$

$$V^*(S3) = 0 + 0.9 (0.9 \times 1000 + 0.1 V^*(S3))$$

$$V^*(S3) = 81000/91.$$

$$V^*(S4) = r(S4, D) + 0.9 (0.9 V^*(S2) + 0.1 V^*(S4))$$

$$V^*(S4) = 40 + 0.9 (0.9 \times 1000 + 0.1 V^*(S4))$$

$$V^*(S4) = 85000/91.$$

# Solve the MDP



Markov Decision Processes satisfy both mentioned properties.

- Bellman equation gives us recursive decomposition (the first property).
- Bellman equation tells us how to break down the optimal value function into two pieces,
  - the optimal behaviour for one step followed by
  - the optimal behaviour after that step.
- We can do prediction, i.e., evaluate the given policy to get the value function on that policy (Dynamic Programming).
- *Evaluating a random policy*
- *Policy Update*

# MDP solving - Value Iteration Example

- In this Grid World, we get a reward of as shown in the figure for each transition we make (actions we take). And the actions that we can take are **north, south, east and west**.
- Our agent is following some random policy  $\pi$  **with a weight of 0.5, 0.3, 0.1 and 0.1** for moving to north, south, east and west directions, respectively.



We need to calculate the policy for four actions, **up, left, right and down** with the center node with a reward of 1



state space and  
reward

# MDP solving – Value Iteration Example


$$\pi_{k+1}(s_i) = \arg \max_a \left\{ r_i + \gamma \sum_{j=1}^N P_{ij}^a V^{\pi_n}(s_j) \right\}$$




state space and  
reward

# MDP solving – Value Iteration Example

$$\pi_{k+1}(s_i) = r_i + \arg \max_a \left\{ \gamma \sum_{j=1}^N P_{ij}^a V^{\pi_n}(s_j) \right\}$$



(left  $\leftarrow$ )  $0.5 * 5 + 0.1 * 10 + 0.1 * 1 + 0.3 * -8 = 1.2$   
 (up  $\uparrow$ )  $0.5 * 10 + 0.1 * 5 + 0.1 * -8 + 0.3 * 1 = 5.0$   
 (right  $\rightarrow$ )  $0.5 * -8 + 0.1 * 10 + 0.1 * 1 + 0.3 * 5 = -1.4$   
 (Down  $\downarrow$ )  $0.5 * 1 + 0.1 * 5 + 0.1 * -8 + 0.3 * 10 = 3.2$



state space and  
reward

$$\pi_1 = 1 + \max\{1.2, 5.0, -1.4, 3.2\} = 6.0 \quad \uparrow$$

# Planning and Approximate Reasoning

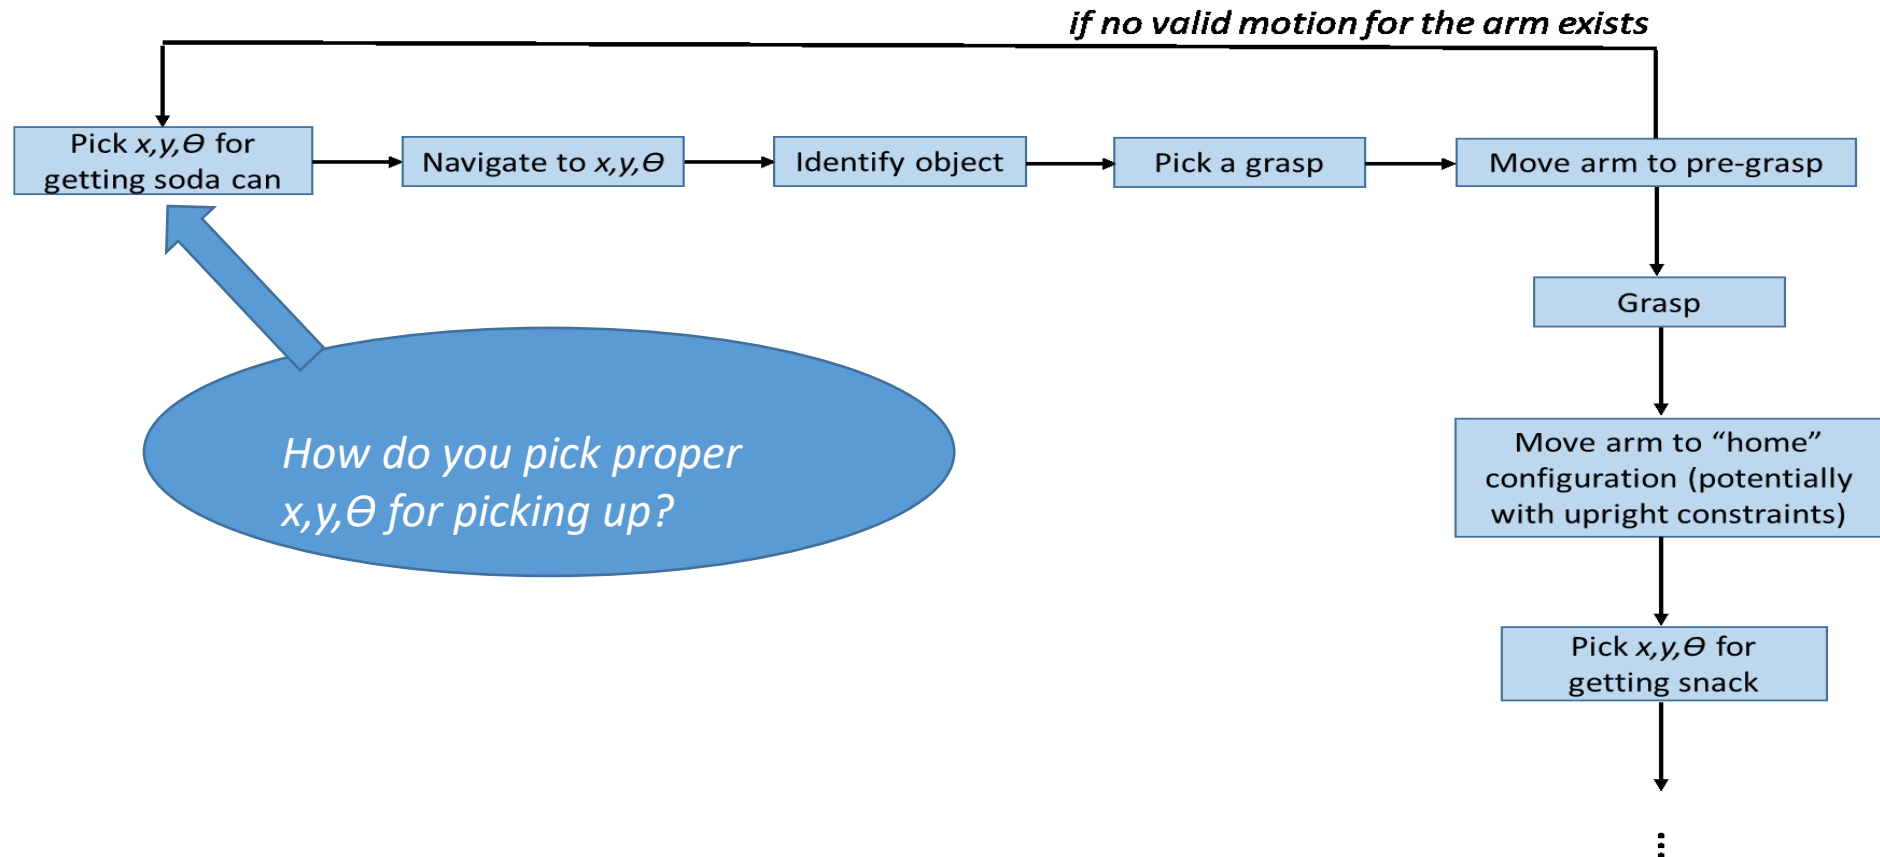
Hatem A. Rashwan

## **Application: Planning for Mobile Robot Manipulation**



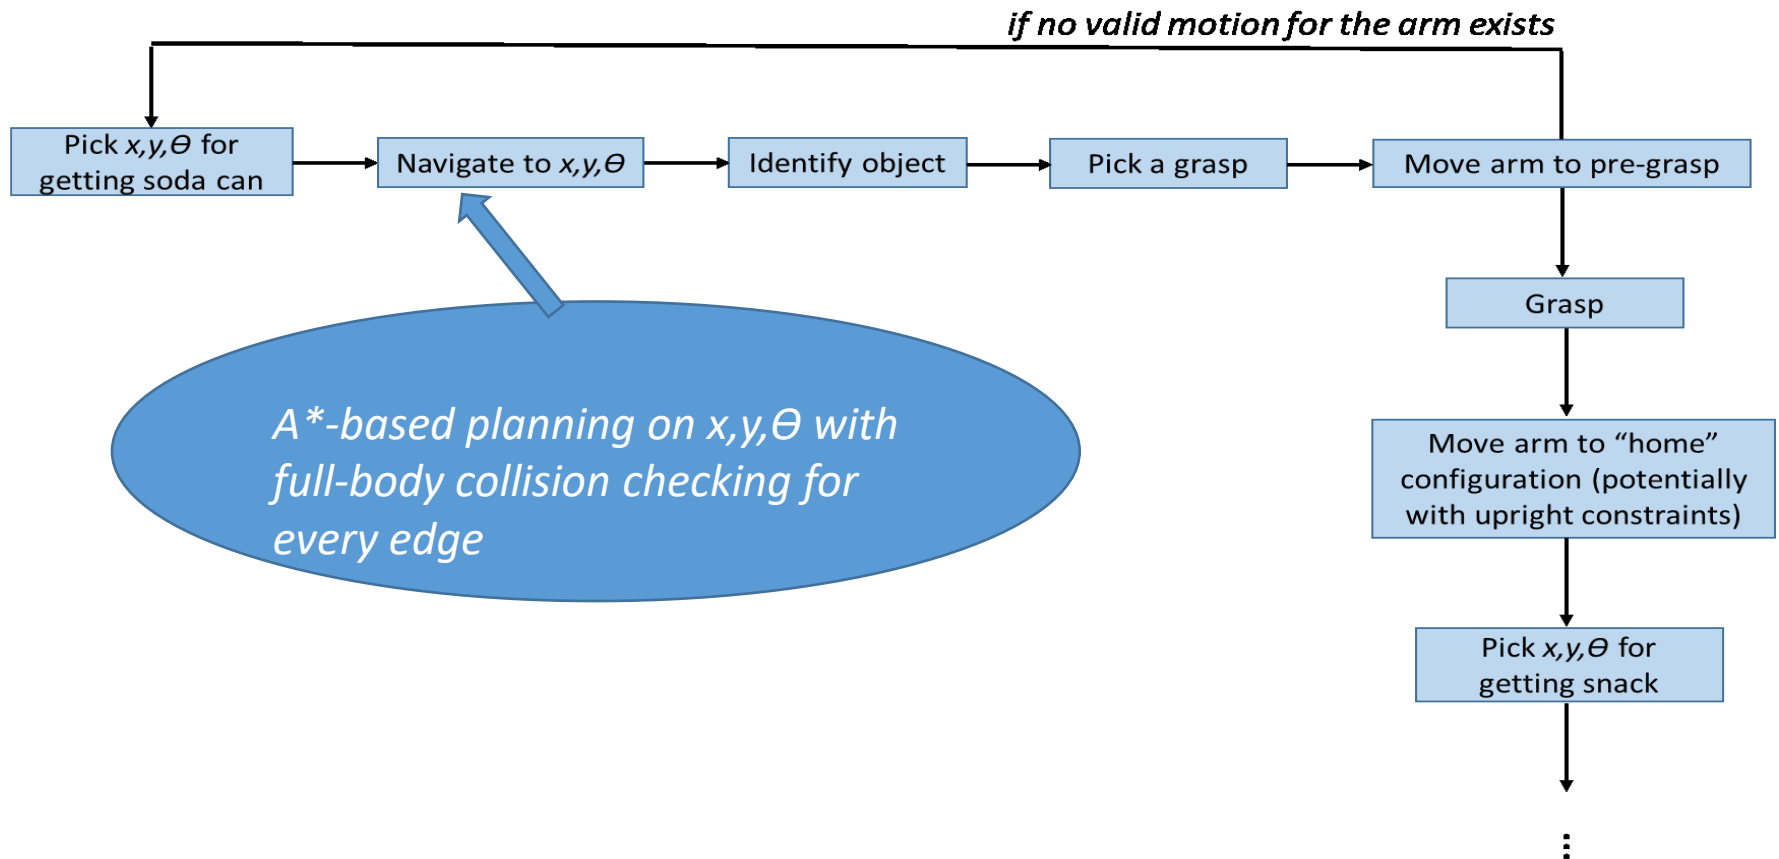
# Robotic Bartender Demo ([Phillips et al.])

- Robot takes in a command from User Interface as to what soda can and snack to deliver



# Robotic Bartender Demo ([Phillips et al.])


- Robot takes in a command from User Interface as to what soda can and snack to deliver



# Graph for Navigation with Complex 3D Body

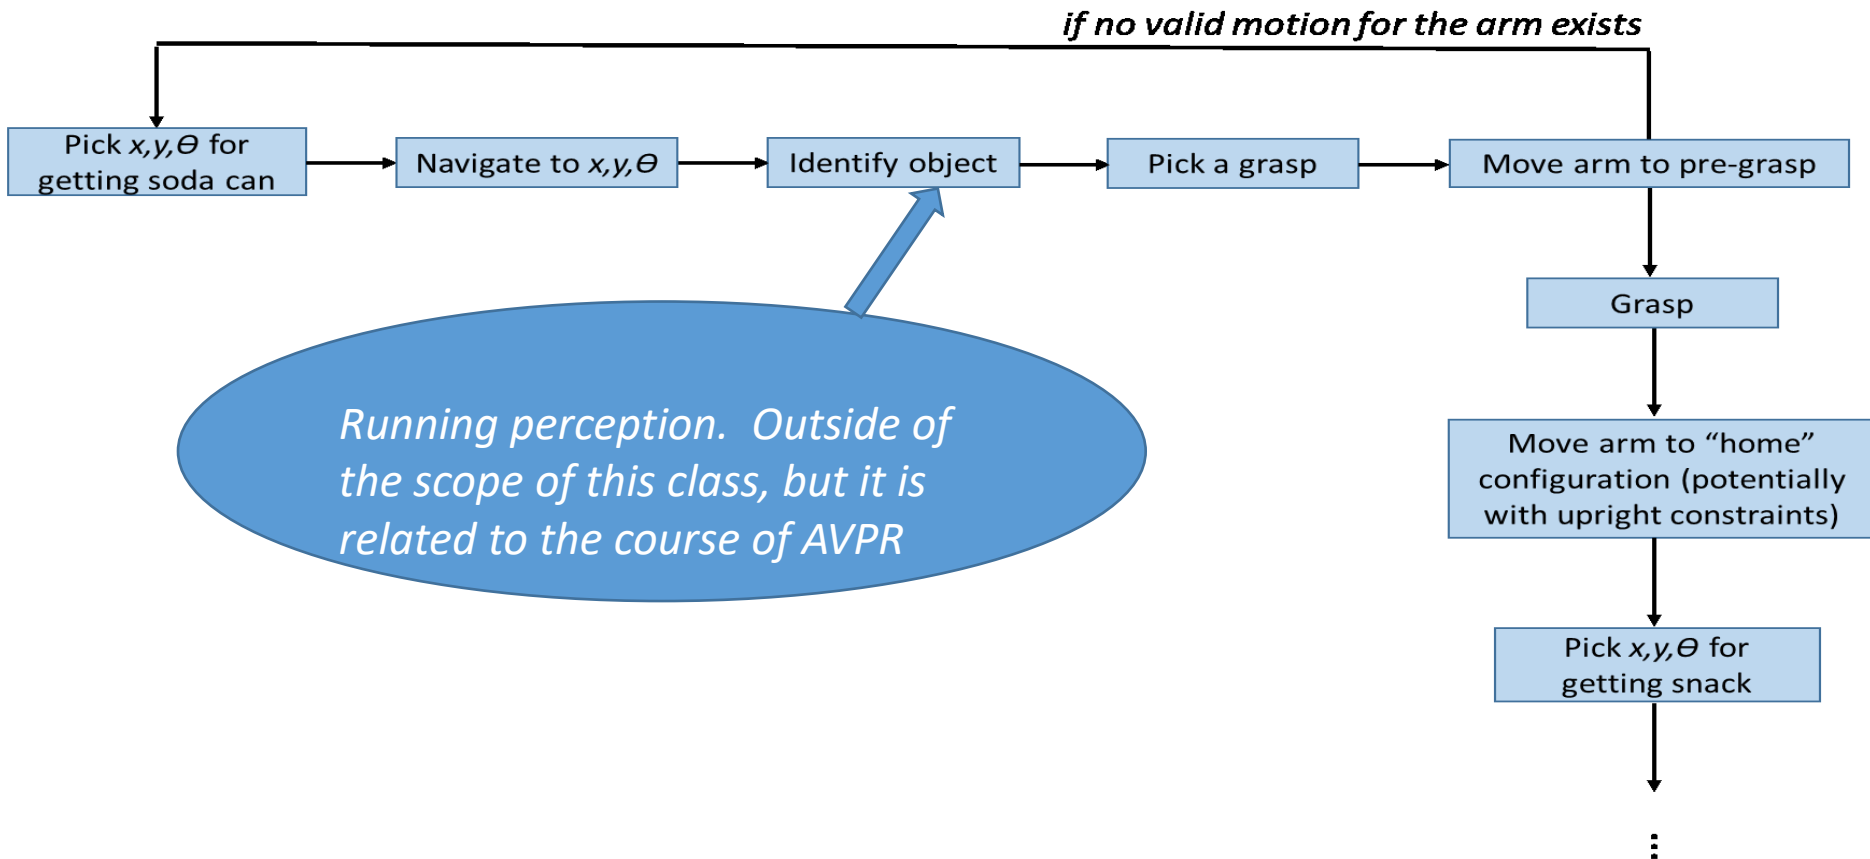
[Hornung et al., '12]

- 3D  $(x, y, \theta)$  lattice-based graph representation for full-body collision checking
  - takes set of motion primitives as input
  - takes  $N$  footprints of the robot defined as polygons as input
  - each footprint corresponds to the projection of a part of the body onto  $x, y$  plane
  - collision checking/cost computation is done for each footprint at the corresponding projection of the 3D map



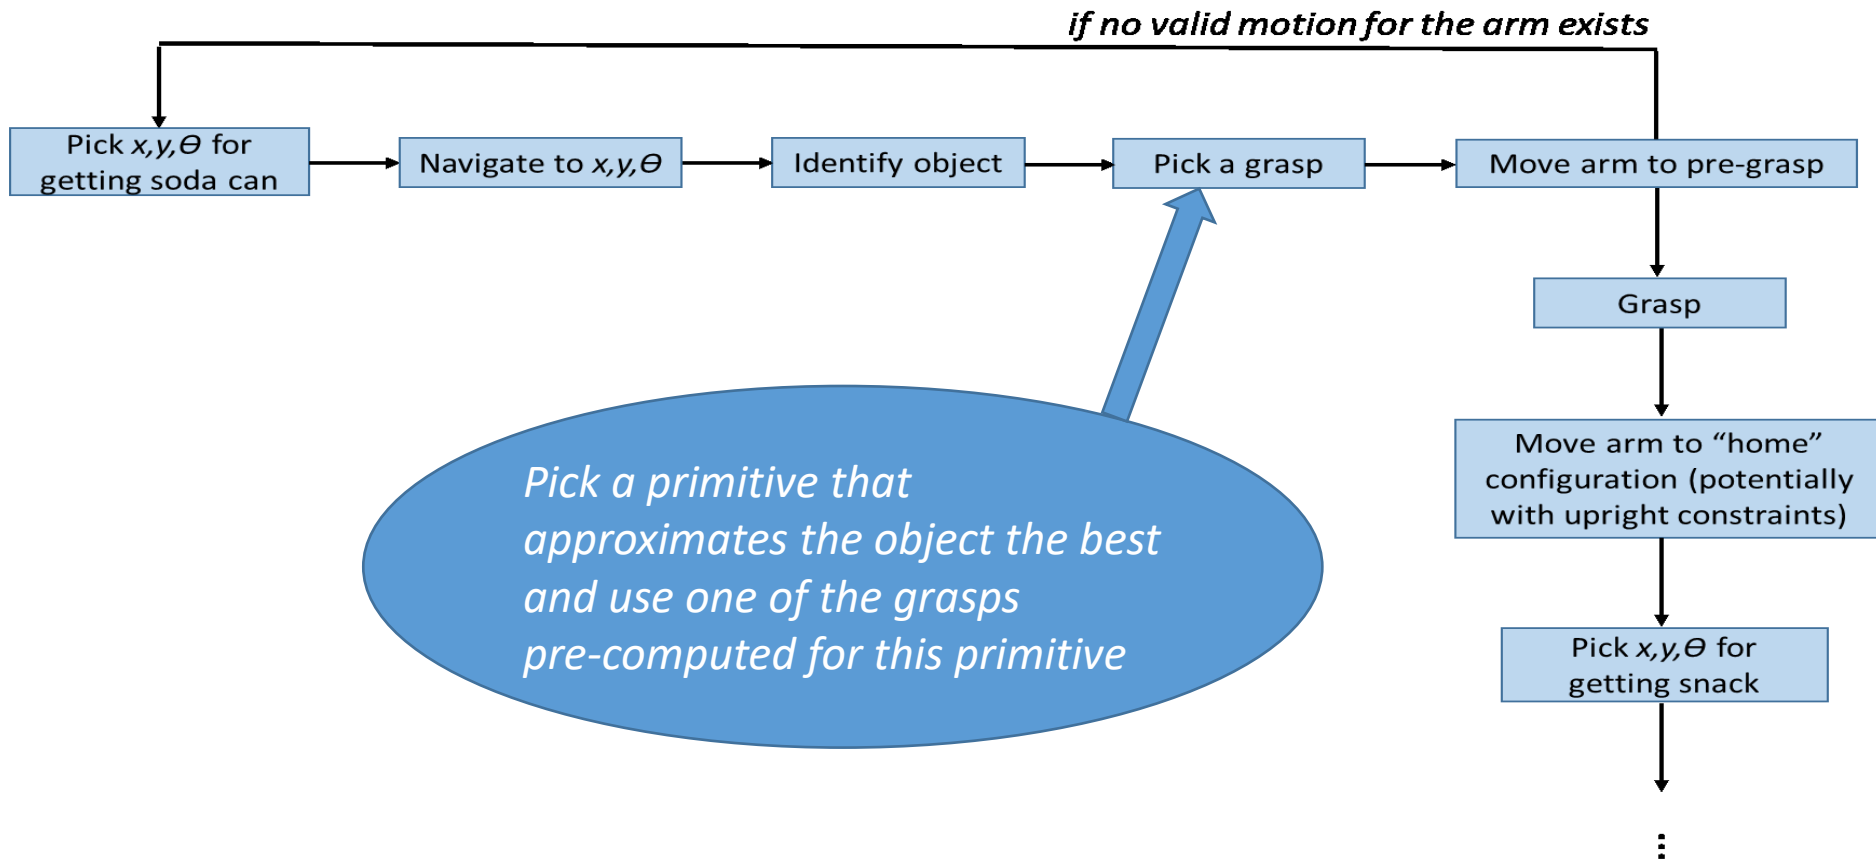
# Robotic Bartender Demo ([Phillips et al.])

- Robot takes in a command from User Interface as to what soda can and snack to deliver



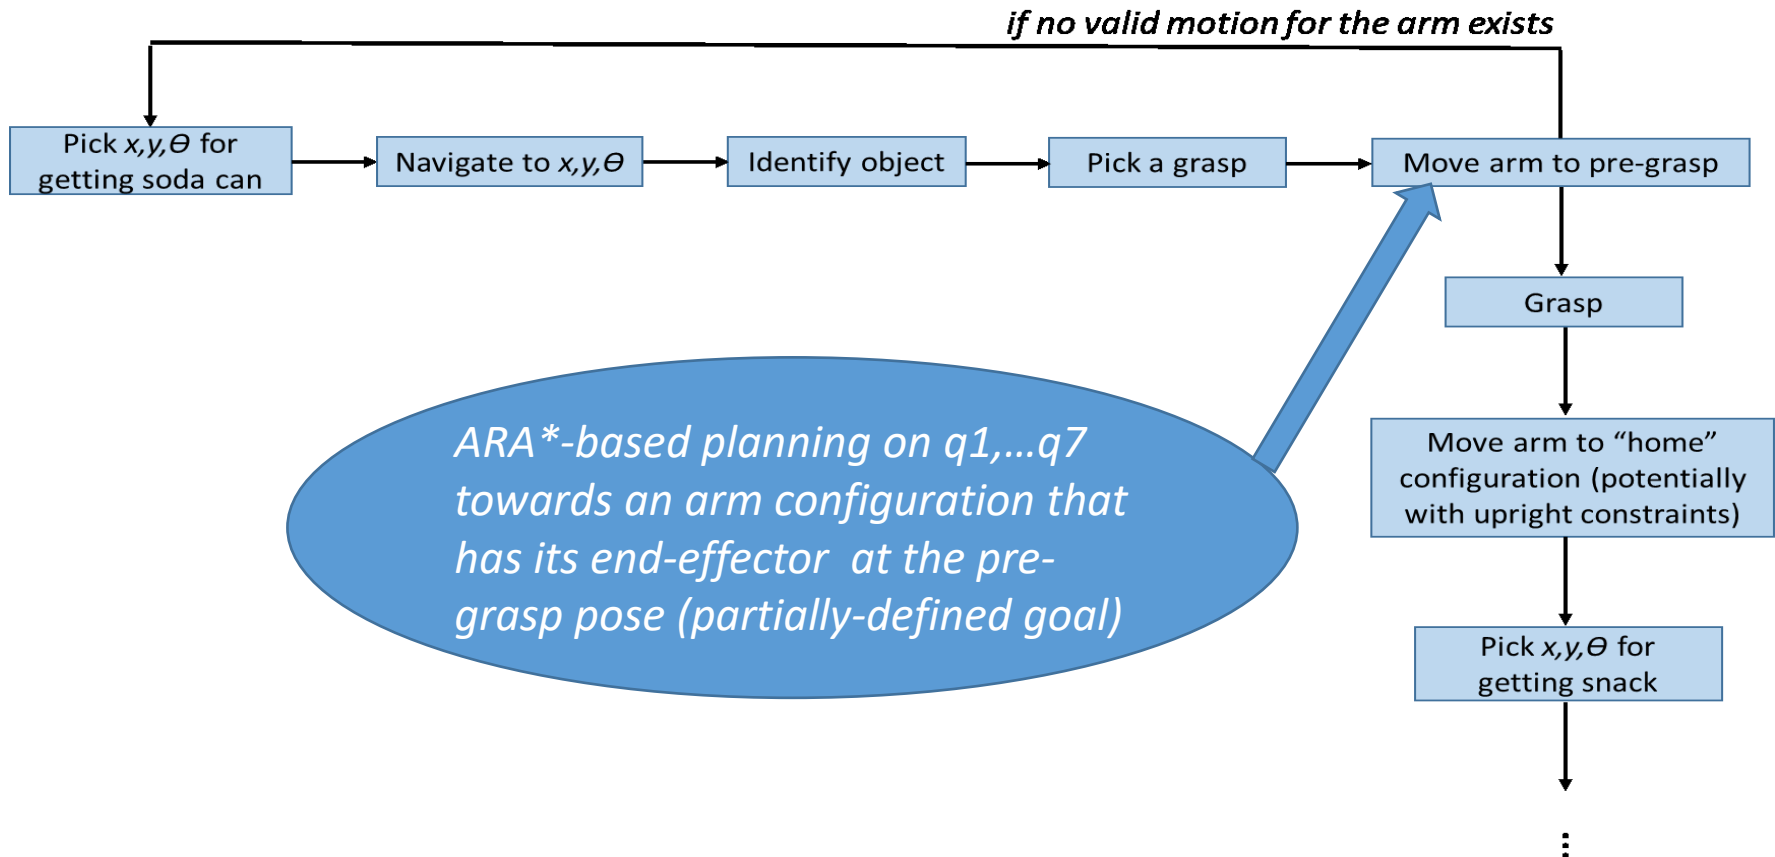
# Robotic Bartender Demo ([Phillips et al.])

- Robot takes in a command from User Interface as to what soda can and snack to deliver



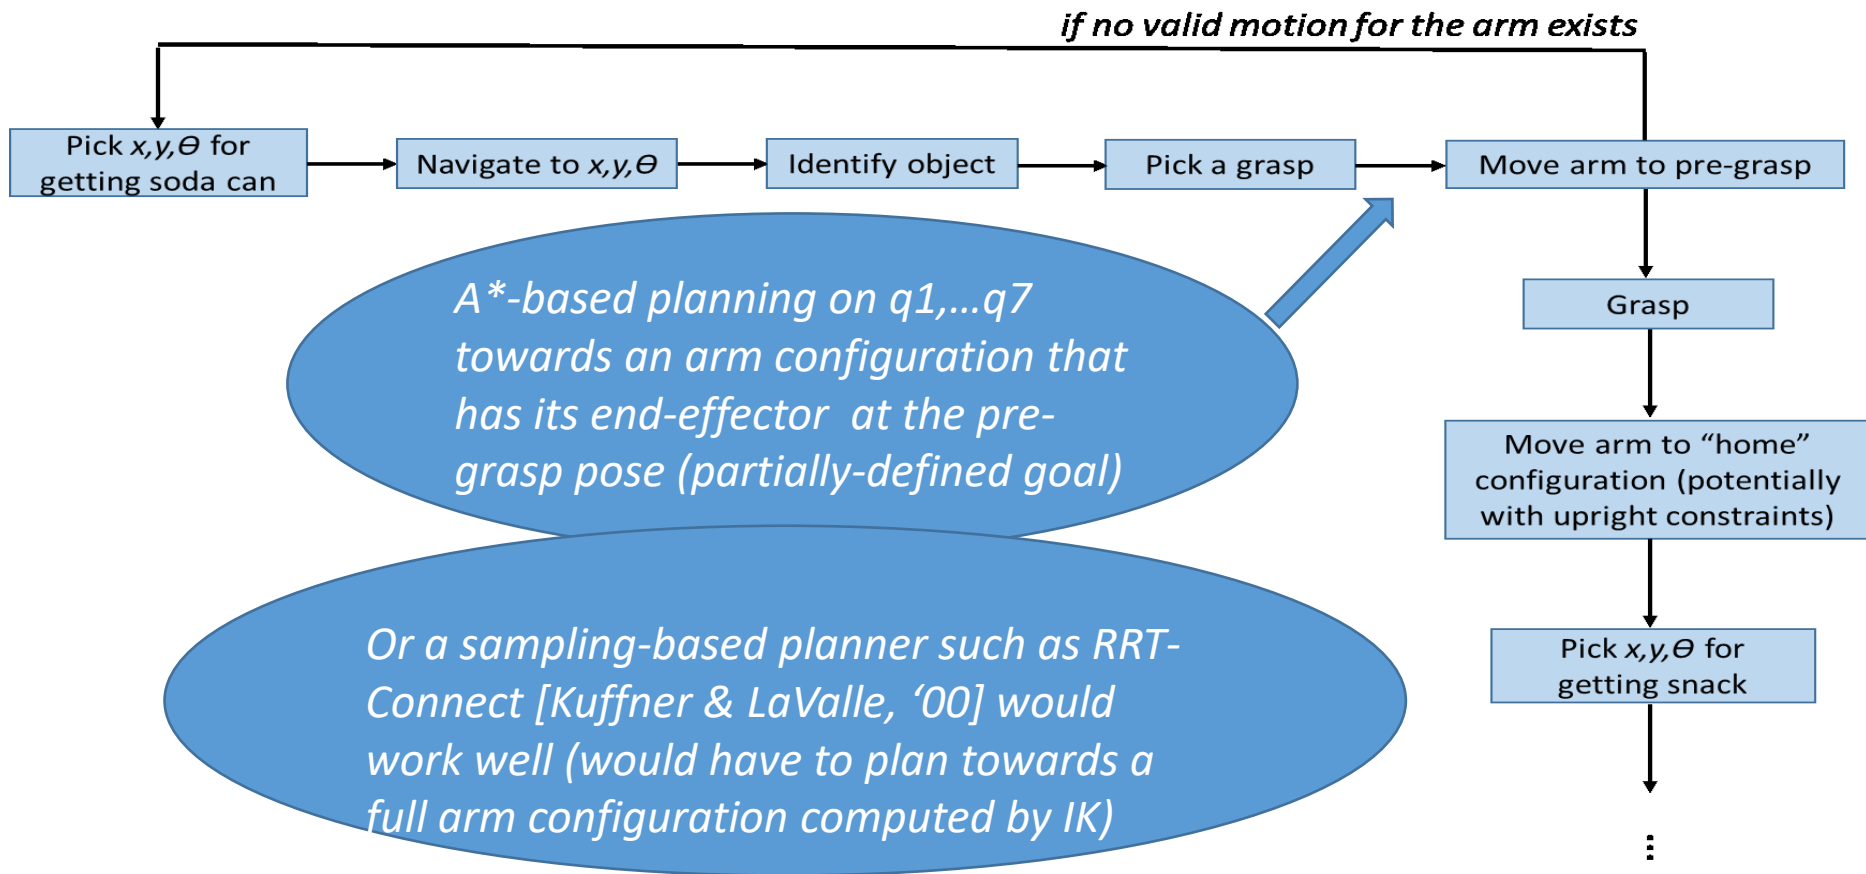
# Robotic Bartender Demo ([Phillips et al.])

- Robot takes in a command from User Interface as to what soda can and snack to deliver



# Robotic Bartender Demo ([Phillips et al.])

- Robot takes in a command from User Interface as to what soda can and snack to deliver

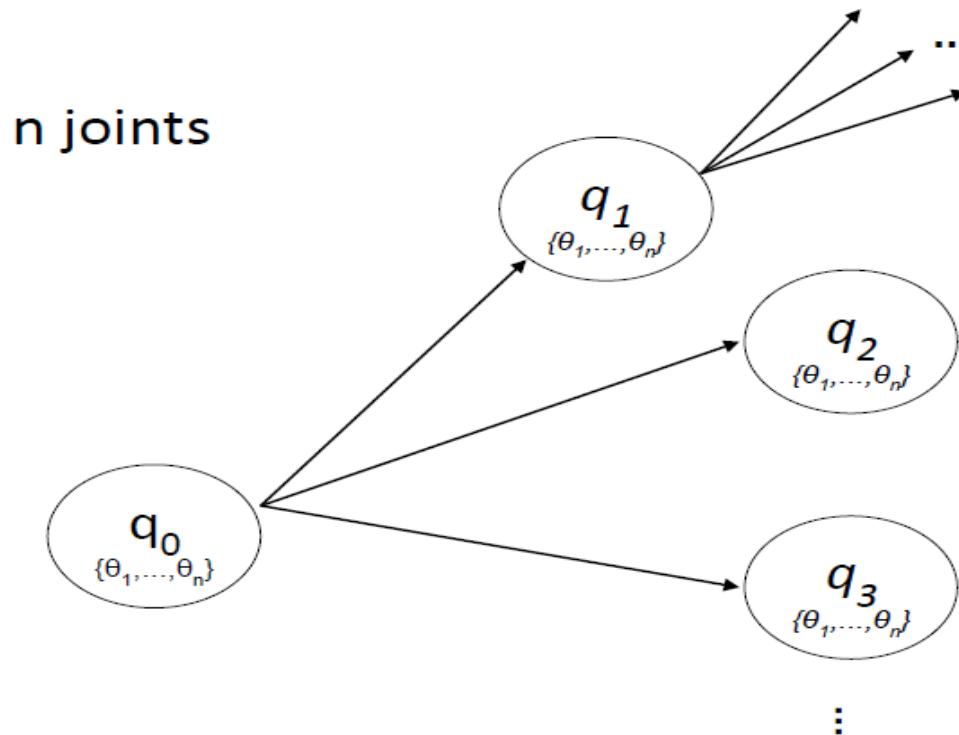


# Manipulation Lattice for Arm Planning

[Cohen et al., '13]

- Representation

- ex. Single arm with  $n$  joints  $\{\theta_1, \dots, \theta_n\}$






# Manipulation Lattice for Arm Planning

[Cohen et al., '13]


- Representation
  - ex. Single arm with  $n$  joints  $\{\theta_1, \dots, \theta_n\}$
- Motion Primitives



# Manipulation Lattice for Arm Planning

## [Cohen et al., '13]

- Non-uniform Dimensionality
  - far from goal: only 4 DoFs active
  - around goal: all 7 DoFs are active
- Non-uniform Resolution
  - far from goal: larger discretization of joint angles
  - around goal: finer discretization of joint angles



# Summary

- ❑ Multiple planners used for both domains
- ❑ Start and goal configurations are often most constrained
  - can be exploited by the planners
- ❑ Planning is higher-dimensional but can take longer than on ground and aerial vehicles
- ❑ Design of proper heuristics is a key to efficiency

# End

